

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.nn.utils.rnn import pad_sequence
from torch.utils.data import TensorDataset, DataLoader, RandomSampler, SequentialSampler
from torchtext import datasets

import numpy as np
from sklearn.metrics import classification_report

import random

from gensim.models import FastText
from nltk import word_tokenize
from nltk.stem import PorterStemmer
from sklearn.metrics import accuracy_score, f1_score

SEED = 1234

random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

train_data, _, test_data = datasets.UDPOS()
train_data = [d for d in train_data]
test_data = [d for d in test_data]

train_tokens = [ [w.lower() for w in d[0]] for d in train_data]
train_tags = [ d[1] for d in train_data]

test_tokens = [[w.lower() for w in d[0]] for d in test_data]
test_tags = [d[1] for d in test_data]

tag2num = { t:i for i, t in enumerate(np.unique([tag for tags in train_tags for tag in tag

100%|██████████| 688k/688k [00:00<00:00, 16.4MB/s]

print(train_data[0][0])
print(train_data[0][1])

print(train_tokens[0])
print(train_tags[0])

['Al', '-', 'Zaman', ':', 'American', 'forces', 'killed', 'Shaikh', 'Abdullah', 'al',
['PROPN', 'PUNCT', 'PROPN', 'PUNCT', 'ADJ', 'NOUN', 'VERB', 'PROPN', 'PROPN', 'PROPN
['al', '-', 'zaman', ':', 'american', 'forces', 'killed', 'shaikh', 'abdullah', 'al',
['PROPN', 'PUNCT', 'PROPN', 'PUNCT', 'ADJ', 'NOUN', 'VERB', 'PROPN', 'PROPN', 'PROPN

```



```

ft = FastText(sentences=train_tokens, size=100, window=5, min_count=1, min_n=1, workers=4)

```

```

max_len = 20
pad_inds = len(tag2num)

def prepare_data(all_tokens, all_tags, ft, tag2num, max_len, pad_tags):
    '''
        Из массива слов all_tokens получим тензор векторов, где каждое слово представлено вект
        А целевую переменную классов all_tags преобразуем в числа.
        Все строки не длинее max_len.
        Пустые значения заполняются нулями или pad_tags
    '''

    # укорачиваем токены
    all_tokens = [tokens[:max_len] for tokens in all_tokens]
    all_tags = [tags[:max_len] for tags in all_tags]

    # переводим теги в числа
    all_tags = [np.array([tag2num[tag] for tag in tags]) for tags in all_tags]

    # all_ids = []
    # for tokens in all_tokens:
    #     ids = prepare_sequence(tokens, word_to_ix)
    #     all_ids.append(ids)

    X_vecs = []
    Y_vecs = []

    for tokens, tags in zip(all_tokens, all_tags):
        X_vecs.append(torch.tensor(np.row_stack([ft.wv[w] for w in tokens])))
        Y_vecs.append(torch.tensor(tags, dtype=torch.long))

    # в качестве заполнителя X используем новый индекс len(word_to_ix)
    X = pad_sequence(X_vecs, batch_first=True)

    # в качестве заполнителя Y используем pad_tags
    Y = pad_sequence(Y_vecs, batch_first=True, padding_value=pad_tags)

    return X, Y

X_train, Y_train = prepare_data(train_tokens, train_tags, ft, tag2num, max_len, pad_inds)

# X_train.size(), Y_train.size()

X_test, Y_test = prepare_data(test_tokens, test_tags, ft, tag2num, max_len, pad_inds)

# X_test.size(), Y_test.size()

X_train.size(), Y_train.size()

(torch.Size([12543, 20, 100]), torch.Size([12543, 20]))

X_test.size(), Y_test.size()

```

```
(torch.Size([2077, 20, 100]), torch.Size([2077, 20]))
```

```
from torch.utils.data import TensorDataset, DataLoader, RandomSampler, SequentialSampler
```

```
bs = 128
```

```
data = TensorDataset(X_train, Y_train)
```

```
dataloader = DataLoader(data, sampler=SequentialSampler(data), batch_size=bs)
```

```
class BiLSTMPOSTagger(nn.Module):
```

```
    def __init__(self, input_dim, hidden_dim, output_dim, n_layers, bidirectional, dropout
```

```
        super().__init__()
```

```
        self.output_dim = output_dim
```

```
        self.input_size = input_dim
```

```
        self.lstm = nn.LSTM(input_dim, hidden_dim, num_layers=n_layers, bidirectional=bidi
```

```
        self.fc = nn.Linear(hidden_dim * 2 if bidirectional else hidden_dim, output_dim)
```

```
        self.dropout = nn.Dropout(dropout)
```

```
    def forward(self, sentence):
```

```
        # sentence = [batch size, sent len, emb dim]
```

```
        sentence = sentence.view(sentence.shape[1], sentence.shape[0], self.input_size)
```

```
        # sentence = [sent len, batch size, emb dim]
```

```
        outputs, (hidden, cell) = self.lstm(sentence)
```

```
        predictions = self.fc(self.dropout(outputs))
```

```
        # predictions = [sent len, batch size, output dim]
```

```
        predictions = predictions.view(predictions.shape[1], predictions.shape[0], self.outpu
```

```
        # predictions = [batch size, sent len, output dim]
```

```
        # raise NotImplementedError()
```

```
        return predictions
```

```
def train_on_epoch(model, dataloader, optimizer):
```

```
    model.train()
```

```
    for batch in dataloader:
```

```
        batch = tuple(t.to(device) for t in batch)
```

```
        b_input, b_tags = batch
```

```
        model.zero_grad()
```

```
        outputs = model(b_input)
```

```
        # outputs = [batch size, sent len, out dim]
```

```
        outputs = outputs.view(-1, outputs.shape[-1])
```

```
        # outputs = [batch size * sent len, out dim]
```

```
        # b_tags = [batch size, sent len]
```

```

        b_tags = b_tags.view(-1)
        # b_tags = [batch size * sent len]

        loss = criterion(outputs, b_tags)
        loss.backward()
        optimizer.step()

def predict_on_dataloader(model, dataloader):
    model.eval()

    all_outputs = []
    all_tags = []
    for batch in dataloader:
        batch = tuple(t.to(device) for t in batch)
        b_input, b_tags = batch
        outputs = model(b_input)

        outputs = outputs.view(-1, outputs.shape[-1])
        b_tags = b_tags.view(-1)

        all_outputs.append(outputs)
        all_tags.append(b_tags)

    all_outputs = torch.cat(all_outputs)
    all_tags = torch.cat(all_tags)

    return all_outputs, all_tags

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)

    cuda

INPUT_DIM = 100
HIDDEN_DIM = 128
OUTPUT_DIM = len(tag2num)
N_LAYERS = 2
BIDIRECTIONAL = True
DROPOUT = 0.25

model = BiLSTMPOSTagger(INPUT_DIM, HIDDEN_DIM, OUTPUT_DIM, N_LAYERS, BIDIRECTIONAL, DROPOUT)
model.to(device)

criterion = nn.CrossEntropyLoss(ignore_index=pad_inds)
optimizer = optim.Adam(model.parameters())

epochs = 50
for e in range(epochs):
    train_on_epoch(model, dataloader, optimizer)

    all_outputs, all_tags = predict_on_dataloader(model, dataloader)
    loss = criterion(all_outputs, all_tags).item()

```

```

all_outputs = all_outputs.detach().cpu().numpy()
all_tags = all_tags.detach().cpu().numpy()

mask = all_tags != pad_inds
loss = loss/len(all_tags[mask])
all_tags = all_tags[mask]
all_preds = np.argmax(all_outputs, axis=1)[mask]

print(f"{e}:\tLoss {loss}, "
      f"accuracy: {accuracy_score(all_tags, all_preds)}, "
      f"f1-macro: {f1_score(all_tags, all_preds, average='macro')}")

```

0:	Loss 1.0015787127812761e-05, accuracy: 0.48026716360080446, f1-macro: 0.251215
1:	Loss 8.605620787400416e-06, accuracy: 0.5540446625706501, f1-macro: 0.3588759
2:	Loss 7.738153386824405e-06, accuracy: 0.5951462819117204, f1-macro: 0.4304601
3:	Loss 7.077609564454969e-06, accuracy: 0.6303437332792118, f1-macro: 0.4783707
4:	Loss 6.623119794917055e-06, accuracy: 0.6454670135365351, f1-macro: 0.4937901
5:	Loss 6.311222655833707e-06, accuracy: 0.6566787825113625, f1-macro: 0.5072632
6:	Loss 6.073827198705909e-06, accuracy: 0.6675399915127586, f1-macro: 0.5193816
7:	Loss 5.869439980562564e-06, accuracy: 0.677866135291549, f1-macro: 0.52703248
8:	Loss 5.71169847315859e-06, accuracy: 0.6825894696704121, f1-macro: 0.52991742
9:	Loss 5.58071466630315e-06, accuracy: 0.686242673604064, f1-macro: 0.533624975
10:	Loss 5.464418580021986e-06, accuracy: 0.6914580219807254, f1-macro: 0.5401216
11:	Loss 5.332528956471036e-06, accuracy: 0.6975466952034786, f1-macro: 0.5567457
12:	Loss 5.226545837043296e-06, accuracy: 0.7033832112523601, f1-macro: 0.5792283
13:	Loss 5.129611353516807e-06, accuracy: 0.7082541498305627, f1-macro: 0.5892182
14:	Loss 5.044750515754521e-06, accuracy: 0.7131312385837377, f1-macro: 0.5985769
15:	Loss 4.964695895189664e-06, accuracy: 0.715861916271518, f1-macro: 0.60115755
16:	Loss 4.8954472902258876e-06, accuracy: 0.7185064915096835, f1-macro: 0.606478
17:	Loss 4.810539163773307e-06, accuracy: 0.7227070610158859, f1-macro: 0.6159342
18:	Loss 4.74201711837671e-06, accuracy: 0.7255299913282532, f1-macro: 0.62670118
19:	Loss 4.681432974527906e-06, accuracy: 0.7288264851135015, f1-macro: 0.6392526
20:	Loss 4.626134166622395e-06, accuracy: 0.7313234561523275, f1-macro: 0.6510896
21:	Loss 4.56079439347741e-06, accuracy: 0.7354502235588602, f1-macro: 0.66328558
22:	Loss 4.5238402981957155e-06, accuracy: 0.7357392817825668, f1-macro: 0.662316
23:	Loss 4.481206427972602e-06, accuracy: 0.7373198767504936, f1-macro: 0.6632319
24:	Loss 4.406753134444333e-06, accuracy: 0.741520446256696, f1-macro: 0.67154366
25:	Loss 4.363998293153021e-06, accuracy: 0.7431071913995954, f1-macro: 0.6741347
26:	Loss 4.322358585163103e-06, accuracy: 0.7458378690873755, f1-macro: 0.6771111
27:	Loss 4.271823837870648e-06, accuracy: 0.7477997749035961, f1-macro: 0.6807209
28:	Loss 4.240151046163306e-06, accuracy: 0.7494787726710825, f1-macro: 0.6829086
29:	Loss 4.18690874641717e-06, accuracy: 0.7520925970343856, f1-macro: 0.68718359
30:	Loss 4.135663569033294e-06, accuracy: 0.754915527346753, f1-macro: 0.69284251
31:	Loss 4.092793988516691e-06, accuracy: 0.75714189068679, f1-macro: 0.695745535
32:	Loss 4.066884918340519e-06, accuracy: 0.7585379804055425, f1-macro: 0.7000836
33:	Loss 4.026599353156436e-06, accuracy: 0.7608627465451392, f1-macro: 0.7020745
34:	Loss 3.976631703568535e-06, accuracy: 0.7633043660092129, f1-macro: 0.7055223
35:	Loss 3.945131937146617e-06, accuracy: 0.7652724220004059, f1-macro: 0.7072132
36:	Loss 3.896189242427453e-06, accuracy: 0.7676709902396723, f1-macro: 0.7138451
37:	Loss 3.862625269687275e-06, accuracy: 0.7697928006051772, f1-macro: 0.7148866
38:	Loss 3.8197406594318956e-06, accuracy: 0.7727510347669391, f1-macro: 0.720878
39:	Loss 3.7906287885804183e-06, accuracy: 0.7744423328843706, f1-macro: 0.722762
40:	Loss 3.7456737401638717e-06, accuracy: 0.7760782794270497, f1-macro: 0.724956
41:	Loss 3.713377764165811e-06, accuracy: 0.779350172512408, f1-macro: 0.72937697
42:	Loss 3.6746670223422764e-06, accuracy: 0.7805433064570687, f1-macro: 0.732772
43:	Loss 3.640653257175796e-06, accuracy: 0.7825728641979864, f1-macro: 0.7343082
44:	Loss 3.607320962360168e-06, accuracy: 0.7839812542666839, f1-macro: 0.7403553
45:	Loss 3.568171058901173e-06, accuracy: 0.7859677607827943, f1-macro: 0.7427956
46:	Loss 3.5515547662366166e-06, accuracy: 0.7877390111748679, f1-macro: 0.742866

```

47:      Loss 3.5106302538650488e-06, accuracy: 0.7896209647164462, f1-macro: 0.747546
48:      Loss 3.496188874375104e-06, accuracy: 0.7896148145414738, f1-macro: 0.7468186
49:      Loss 3.4430745106980625e-06, accuracy: 0.7930343118261715, f1-macro: 0.752476

```

```

def count_metrics(model, dataloader):
    y_pred, y_true = predict_on_dataloader(model, dataloader)

    y_pred = y_pred.detach().cpu().numpy()
    y_true = y_true.detach().cpu().numpy()

    mask = y_true != pad_inds
    y_true = y_true[mask]
    y_pred = np.argmax(y_pred, axis=1)[mask]

    print(classification_report(y_true, y_pred))

```

```
count_metrics(model, dataloader)
```

	precision	recall	f1-score	support
0	0.63	0.36	0.46	9962
1	0.89	0.85	0.87	13578
2	0.70	0.70	0.70	8547
3	0.86	0.95	0.90	10404
4	0.99	0.99	0.99	5202
5	0.95	0.98	0.96	13014
6	0.95	0.63	0.76	649
7	0.58	0.86	0.69	27080
8	0.94	0.94	0.94	3339
9	0.76	0.93	0.83	4484
10	0.95	0.95	0.95	15619
11	0.76	0.17	0.28	10523
12	0.99	1.00	0.99	16990
13	0.77	0.60	0.67	3134
14	0.86	0.78	0.82	484
15	0.74	0.70	0.72	18849
16	0.96	0.14	0.25	739
accuracy			0.79	162597
macro avg	0.84	0.74	0.75	162597
weighted avg	0.80	0.79	0.78	162597

```

data = TensorDataset(X_test, Y_test)
test_dataloader = DataLoader(data, sampler=SequentialSampler(data), batch_size=bs)
count_metrics(model, test_dataloader)

```


	precision	recall	f1-score	support
0	0.59	0.35	0.44	1466
1	0.86	0.83	0.84	1656
2	0.64	0.64	0.64	1066
3	0.81	0.91	0.86	1336
4	0.99	0.99	0.99	599
5	0.94	0.97	0.96	1607

6	0.94	0.43	0.59	115
7	0.51	0.80	0.63	3446
8	0.91	0.95	0.93	448
9	0.70	0.83	0.76	546
10	0.92	0.93	0.92	1923
11	0.62	0.08	0.15	1773
12	0.98	0.99	0.98	2467
13	0.62	0.45	0.52	330
14	0.67	0.56	0.61	81
15	0.64	0.68	0.66	2306
16	0.00	0.00	0.00	114
accuracy				0.74 21279
macro avg				0.73 0.67 0.67 21279
weighted avg				0.74 0.74 0.72 21279

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
_warn_prf(average, modifier, msg_start, len(result))

```



✓ 0 сек. выполнено в 14:55

